

# Data Analysis and Machine Learning: Linear Regression

Morten Hjorth-Jensen<sup>1,2</sup>

<sup>1</sup>Department of Physics and Center for Computing in Science Education, University of Oslo, Norway

<sup>2</sup>Department of Physics and Astronomy and Facility for Rare Ion Beams and National Superconducting Cyclotron Laboratory, Michigan State University, U

Jun 23, 2020

## Linear Regression, basic overview

The aim of this set of lectures is to introduce basic aspects of linear regression, a widely applied set of methods used to fit continuous functions.

We will also use these widely popular methods to introduce resampling techniques like bootstrapping and cross-validation.

We will in particular focus on

- Ordinary linear regression
- Ridge regression
- Lasso regression
- Resampling techniques
- Bias-variance tradeoff

## Why Linear Regression (aka Ordinary Least Squares and family)?

Fitting a continuous function with linear parameterization in terms of the parameters  $\beta$ .

- Method of choice for fitting a continuous function!
- Gives an excellent introduction to central Machine Learning features with **understandable pedagogical** links to other methods like **Neural Networks, Support Vector Machines** etc
- Analytical expression for the fitting parameters  $\beta$

- Analytical expressions for statistical properties like mean values, variances, confidence intervals and more
- Analytical relation with probabilistic interpretations
- Easy to introduce basic concepts like bias-variance tradeoff, cross-validation, resampling and regularization techniques and many other ML topics
- Easy to code! And links well with classification problems and logistic regression and neural networks
- Allows for **easy** hands-on understanding of gradient descent methods. These methods are at the heart of all essentially all Machine Learning methods.
- and many more features

## Additional Reading

For more discussions of Ridge and Lasso regression, [Wessel van Wieringen's](#) article is highly recommended. Similarly, [Mehta et al's](#) article is also recommended. The textbook by [Hastie, Tibshirani, and Friedman on The Elements of Statistical Learning Data Mining](#) is highly recommended.

## Regression Analysis, Definitions and Aims

### Regression analysis, overarching aims

Regression modeling deals with the description of the sampling distribution of a given random variable  $y$  and how it varies as function of another variable or a set of such variables  $\mathbf{x} = [x_0, x_1, \dots, x_{n-1}]^T$ . The first variable is called the **dependent**, the **outcome** or the **response** variable while the set of variables  $\mathbf{x}$  is called the independent variable, or the predictor variable or the explanatory variable.

A regression model aims at finding a likelihood function  $p(\mathbf{y}|\mathbf{x})$ , that is the conditional distribution for  $\mathbf{y}$  with a given  $\mathbf{x}$ . The estimation of  $p(\mathbf{y}|\mathbf{x})$  is made using a data set with

- $n$  cases  $i = 0, 1, 2, \dots, n - 1$
- Response (target, dependent or outcome) variable  $y_i$  with  $i = 0, 1, 2, \dots, n - 1$
- $p$  so-called explanatory (independent or predictor) variables  $\mathbf{x}_i = [x_{i0}, x_{i1}, \dots, x_{ip-1}]$  with  $i = 0, 1, 2, \dots, n - 1$  and explanatory variables running from 0 to  $p - 1$ . See below for more explicit examples. These variables are also called features or predictors.

The goal of the regression analysis is to extract/exploit relationship between  $\mathbf{y}$  and  $\mathbf{x}$  in or to infer causal dependencies, approximations to the likelihood functions, functional relationships and to make predictions, making fits and many other things.

## Regression analysis, overarching aims II

Consider an experiment in which  $p$  characteristics of  $n$  samples are measured. The data from this experiment, for various explanatory variables  $p$  are normally represented by a matrix  $\mathbf{X}$ .

The matrix  $\mathbf{X}$  is called the *design matrix*. Additional information of the samples is available in the form of  $\mathbf{y}$  (also as above). The variable  $\mathbf{y}$  is generally referred to as the *response variable*. The aim of regression analysis is to explain  $\mathbf{y}$  in terms of  $\mathbf{X}$  through a functional relationship like  $y_i = f(\mathbf{X}_{i,*})$ . When no prior knowledge on the form of  $f(\cdot)$  is available, it is common to assume a linear relationship between  $\mathbf{X}$  and  $\mathbf{y}$ . This assumption gives rise to the *linear regression model* where  $\boldsymbol{\beta} = [\beta_0, \dots, \beta_{p-1}]^T$  are the *regression parameters*.

Linear regression gives us a set of analytical equations for the parameters  $\beta_j$ .

## Examples

In order to understand the relation among the predictors  $p$ , the set of data  $n$  and the target (outcome, output etc)  $\mathbf{y}$ , consider the model we discussed for describing nuclear binding energies.

There we assumed that we could parametrize the data using a polynomial approximation based on the liquid drop model. Assuming

$$BE(A) = a_0 + a_1 A + a_2 A^{2/3} + a_3 A^{-1/3} + a_4 A^{-1},$$

we have five predictors, that is the intercept, the  $A$  dependent term, the  $A^{2/3}$  term and the  $A^{-1/3}$  and  $A^{-1}$  terms. This gives  $p = 0, 1, 2, 3, 4$ . Furthermore we have  $n$  entries for each predictor. It means that our design matrix is an  $n \times p$  matrix  $\mathbf{X}$ .

Here the predictors are based on a model we have made. A popular data set which is widely encountered in ML applications is the so-called [credit card default data from Taiwan](#). The data set contains data on  $n = 30000$  credit card holders with predictors like gender, marital status, age, profession, education, etc. In total there are 24 such predictors or attributes leading to a design matrix of dimensionality  $24 \times 30000$ . This is however a classification problem and we will come back to it when we discuss Logistic Regression.

## General linear models

Before we proceed let us study a case from linear algebra where we aim at fitting a set of data  $\mathbf{y} = [y_0, y_1, \dots, y_{n-1}]$ . We could think of these data as a result of an experiment or a complicated numerical experiment. These data are functions of a series of variables  $\mathbf{x} = [x_0, x_1, \dots, x_{n-1}]$ , that is  $y_i = y(x_i)$  with

$i = 0, 1, 2, \dots, n-1$ . The variables  $x_i$  could represent physical quantities like time, temperature, position etc. We assume that  $y(x)$  is a smooth function.

Since obtaining these data points may not be trivial, we want to use these data to fit a function which can allow us to make predictions for values of  $y$  which are not in the present set. The perhaps simplest approach is to assume we can parametrize our function in terms of a polynomial of degree  $n-1$  with  $n$  points, that is

$$y = y(x) \rightarrow y(x_i) = \tilde{y}_i + \epsilon_i = \sum_{j=0}^{n-1} \beta_j x_i^j + \epsilon_i,$$

where  $\epsilon_i$  is the error in our approximation.

### Rewriting the fitting procedure as a linear algebra problem

For every set of values  $y_i, x_i$  we have thus the corresponding set of equations

$$\begin{aligned} y_0 &= \beta_0 + \beta_1 x_0^1 + \beta_2 x_0^2 + \dots + \beta_{n-1} x_0^{n-1} + \epsilon_0 \\ y_1 &= \beta_0 + \beta_1 x_1^1 + \beta_2 x_1^2 + \dots + \beta_{n-1} x_1^{n-1} + \epsilon_1 \\ y_2 &= \beta_0 + \beta_1 x_2^1 + \beta_2 x_2^2 + \dots + \beta_{n-1} x_2^{n-1} + \epsilon_2 \\ &\dots\dots\dots \\ y_{n-1} &= \beta_0 + \beta_1 x_{n-1}^1 + \beta_2 x_{n-1}^2 + \dots + \beta_{n-1} x_{n-1}^{n-1} + \epsilon_{n-1}. \end{aligned}$$

### Rewriting the fitting procedure as a linear algebra problem, more details

Defining the vectors

$$\mathbf{y} = [y_0, y_1, y_2, \dots, y_{n-1}]^T,$$

and

$$\boldsymbol{\beta} = [\beta_0, \beta_1, \beta_2, \dots, \beta_{n-1}]^T,$$

and

$$\boldsymbol{\epsilon} = [\epsilon_0, \epsilon_1, \epsilon_2, \dots, \epsilon_{n-1}]^T,$$

and the design matrix

$$\mathbf{X} = \begin{bmatrix} 1 & x_0^1 & x_0^2 & \dots & \dots & x_0^{n-1} \\ 1 & x_1^1 & x_1^2 & \dots & \dots & x_1^{n-1} \\ 1 & x_2^1 & x_2^2 & \dots & \dots & x_2^{n-1} \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n-1}^1 & x_{n-1}^2 & \dots & \dots & x_{n-1}^{n-1} \end{bmatrix}$$

we can rewrite our equations as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}.$$

The above design matrix is called a [Vandermonde matrix](#).

## Generalizing the fitting procedure as a linear algebra problem

We are obviously not limited to the above polynomial expansions. We could replace the various powers of  $x$  with elements of Fourier series or instead of  $x_i^j$  we could have  $\cos(jx_i)$  or  $\sin(jx_i)$ , or time series or other orthogonal functions. For every set of values  $y_i, x_i$  we can then generalize the equations to

$$\begin{aligned}
 y_0 &= \beta_0 x_{00} + \beta_1 x_{01} + \beta_2 x_{02} + \cdots + \beta_{n-1} x_{0n-1} + \epsilon_0 \\
 y_1 &= \beta_0 x_{10} + \beta_1 x_{11} + \beta_2 x_{12} + \cdots + \beta_{n-1} x_{1n-1} + \epsilon_1 \\
 y_2 &= \beta_0 x_{20} + \beta_1 x_{21} + \beta_2 x_{22} + \cdots + \beta_{n-1} x_{2n-1} + \epsilon_2 \\
 &\dots\dots\dots \\
 y_i &= \beta_0 x_{i0} + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_{n-1} x_{in-1} + \epsilon_i \\
 &\dots\dots\dots \\
 y_{n-1} &= \beta_0 x_{n-1,0} + \beta_1 x_{n-1,1} + \beta_2 x_{n-1,2} + \cdots + \beta_{n-1} x_{n-1,n-1} + \epsilon_{n-1}.
 \end{aligned}$$

Note that we have used  $p = n$  here. The matrix is thus quadratic (it may be symmetric). This is generally not the case!

## Generalizing the fitting procedure as a linear algebra problem

We redefine in turn the matrix  $\mathbf{X}$  as

$$\mathbf{X} = \begin{bmatrix} x_{00} & x_{01} & x_{02} & \dots & \dots & x_{0,n-1} \\ x_{10} & x_{11} & x_{12} & \dots & \dots & x_{1,n-1} \\ x_{20} & x_{21} & x_{22} & \dots & \dots & x_{2,n-1} \\ \dots & \dots & \dots & \dots & \dots & \dots \\ x_{n-1,0} & x_{n-1,1} & x_{n-1,2} & \dots & \dots & x_{n-1,n-1} \end{bmatrix}$$

and without loss of generality we rewrite again our equations as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}.$$

The left-hand side of this equation is known. Our error vector  $\boldsymbol{\epsilon}$  and the parameter vector  $\boldsymbol{\beta}$  are our unknown quantities. How can we obtain the optimal set of  $\beta_i$  values?

## Optimizing our parameters

We have defined the matrix  $\mathbf{X}$  via the equations

$$\begin{aligned}y_0 &= \beta_0 x_{00} + \beta_1 x_{01} + \beta_2 x_{02} + \cdots + \beta_{n-1} x_{0n-1} + \epsilon_0 \\y_1 &= \beta_0 x_{10} + \beta_1 x_{11} + \beta_2 x_{12} + \cdots + \beta_{n-1} x_{1n-1} + \epsilon_1 \\y_2 &= \beta_0 x_{20} + \beta_1 x_{21} + \beta_2 x_{22} + \cdots + \beta_{n-1} x_{2n-1} + \epsilon_1 \\&\dots\dots\dots \\y_i &= \beta_0 x_{i0} + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_{n-1} x_{in-1} + \epsilon_i \\&\dots\dots\dots \\y_{n-1} &= \beta_0 x_{n-1,0} + \beta_1 x_{n-1,1} + \beta_2 x_{n-1,2} + \cdots + \beta_{n-1} x_{n-1,n-1} + \epsilon_{n-1}.\end{aligned}$$

As we noted above, we stayed with a system with the design matrix  $\mathbf{X} \in \mathbb{R}^{n \times n}$ , that is we have  $p = n$ . For reasons to come later (algorithmic arguments) we will hereafter define our matrix as  $\mathbf{X} \in \mathbb{R}^{n \times p}$ , with the predictors referring to the column numbers and the entries  $n$  being the row elements.

## Our model for the nuclear binding energies

In our [introductory notes](#) we looked at the so-called [liquid drop model](#). Let us remind ourselves about what we did by looking at the code.

We restate the parts of the code we are most interested in.

With  $\beta \in \mathbb{R}^{p \times 1}$ , it means that we will hereafter write our equations for the approximation as

$$\tilde{\mathbf{y}} = \mathbf{X}\beta,$$

throughout these lectures.

## Optimizing our parameters, more details

With the above we use the design matrix to define the approximation  $\tilde{\mathbf{y}}$  via the unknown quantity  $\beta$  as

$$\tilde{\mathbf{y}} = \mathbf{X}\beta,$$

and in order to find the optimal parameters  $\beta_i$  instead of solving the above linear algebra problem, we define a function which gives a measure of the spread between the values  $y_i$  (which represent hopefully the exact values) and the parameterized values  $\tilde{y}_i$ , namely

$$C(\beta) = \frac{1}{n} \sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2 = \frac{1}{n} \left\{ (\mathbf{y} - \tilde{\mathbf{y}})^T (\mathbf{y} - \tilde{\mathbf{y}}) \right\},$$

or using the matrix  $\mathbf{X}$  and in a more compact matrix-vector notation as

$$C(\beta) = \frac{1}{n} \left\{ (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta) \right\}.$$

This function is one possible way to define the so-called cost function.

It is also common to define the function  $C$  as

$$C(\boldsymbol{\beta}) = \frac{1}{2n} \sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2,$$

since when taking the first derivative with respect to the unknown parameters  $\boldsymbol{\beta}$ , the factor of 2 cancels out.

## Interpretations and optimizing our parameters

The function

$$C(\boldsymbol{\beta}) = \frac{1}{n} \left\{ (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right\},$$

can be linked to the variance of the quantity  $y_i$  if we interpret the latter as the mean value. Below we will show that

$$y_i = \langle y_i \rangle = \beta_0 x_{i,0} + \beta_1 x_{i,1} + \beta_2 x_{i,2} + \cdots + \beta_{n-1} x_{i,n-1} + \epsilon_i,$$

where  $\langle y_i \rangle$  is the mean value. Keep in mind also that till now we have treated  $y_i$  as the exact value. Normally, the response (dependent or outcome) variable  $y_i$  the outcome of a numerical experiment or another type of experiment and is thus only an approximation to the true value. It is then always accompanied by an error estimate, often limited to a statistical error estimate given by the standard deviation discussed earlier. In the discussion here we will treat  $y_i$  as our exact value for the response variable.

In order to find the parameters  $\beta_i$  we will then minimize the spread of  $C(\boldsymbol{\beta})$ , that is we are going to solve the problem

$$\min_{\boldsymbol{\beta} \in \mathbb{R}^p} \frac{1}{n} \left\{ (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right\}.$$

In practical terms it means we will require

$$\frac{\partial C(\boldsymbol{\beta})}{\partial \beta_j} = \frac{\partial}{\partial \beta_j} \left[ \frac{1}{n} \sum_{i=0}^{n-1} (y_i - \beta_0 x_{i,0} - \beta_1 x_{i,1} - \beta_2 x_{i,2} - \cdots - \beta_{n-1} x_{i,n-1})^2 \right] = 0,$$

which results in

$$\frac{\partial C(\boldsymbol{\beta})}{\partial \beta_j} = -\frac{2}{n} \left[ \sum_{i=0}^{n-1} x_{ij} (y_i - \beta_0 x_{i,0} - \beta_1 x_{i,1} - \beta_2 x_{i,2} - \cdots - \beta_{n-1} x_{i,n-1}) \right] = 0,$$

or in a matrix-vector form as

$$\frac{\partial C(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = 0 = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}).$$

## Interpretations and optimizing our parameters

We can rewrite

$$\frac{\partial C(\beta)}{\partial \beta} = 0 = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta),$$

as

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X} \beta,$$

and if the matrix  $\mathbf{X}^T \mathbf{X}$  is invertible we have the solution

$$\beta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

We note also that since our design matrix is defined as  $\mathbf{X} \in \mathbb{R}^{n \times p}$ , the product  $\mathbf{X}^T \mathbf{X} \in \mathbb{R}^{p \times p}$ . In the above case we have that  $p \ll n$ , in our case  $p = 5$  meaning that we end up with inverting a small  $5 \times 5$  matrix. This is a rather common situation, in many cases we end up with low-dimensional matrices to invert. The methods discussed here and for many other supervised learning algorithms like classification with logistic regression or support vector machines, exhibit dimensionalities which allow for the usage of direct linear algebra methods such as **LU** decomposition or **Singular Value Decomposition** (SVD) for finding the inverse of the matrix  $\mathbf{X}^T \mathbf{X}$ .

**Small question:** Do you think the example we have at hand here (the nuclear binding energies) can lead to problems in inverting the matrix  $\mathbf{X}^T \mathbf{X}$ ? What kind of problems can we expect?

## Some useful matrix and vector expressions

The following matrix and vector relation will be useful here and for the rest of the course. Vectors are always written as boldfaced lower case letters and matrices as upper case boldfaced letters.

$$\begin{aligned}\frac{\partial(\mathbf{b}^T \mathbf{a})}{\partial \mathbf{a}} &= \mathbf{b}, \\ \frac{\partial(\mathbf{a}^T \mathbf{A} \mathbf{a})}{\partial \mathbf{a}} &= (\mathbf{A} + \mathbf{A}^T) \mathbf{a}, \\ \frac{\partial \text{tr}(\mathbf{B} \mathbf{A})}{\partial \mathbf{A}} &= \mathbf{B}^T, \\ \frac{\partial \log |\mathbf{A}|}{\partial \mathbf{A}} &= (\mathbf{A}^{-1})^T.\end{aligned}$$

## Interpretations and optimizing our parameters

The residuals  $\epsilon$  are in turn given by

$$\epsilon = \mathbf{y} - \tilde{\mathbf{y}} = \mathbf{y} - \mathbf{X}\beta,$$



and with

$$\mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) = 0,$$

we have

$$\mathbf{X}^T \boldsymbol{\epsilon} = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) = 0,$$

meaning that the solution for  $\boldsymbol{\beta}$  is the one which minimizes the residuals. Later we will link this with the maximum likelihood approach.

Let us now return to our nuclear binding energies and simply code the above equations.

## Own code for Ordinary Least Squares

It is rather straightforward to implement the matrix inversion and obtain the parameters  $\boldsymbol{\beta}$ . After having defined the matrix  $\mathbf{X}$  we simply need to write Alternatively, you can use the least squares functionality in **Numpy** as

And finally we plot our fit with and compare with data

## Adding error analysis and training set up

We can easily test our fit by computing the  $R^2$  score that we discussed in connection with the functionality of *scikit\_Learn* in the introductory slides. Since we are not using *scikit-Learn* here we can define our own  $R^2$  function as and we would be using it as

We can also add our **MSE** score as and finally the relative error as

## The $\chi^2$ function

Normally, the response (dependent or outcome) variable  $y_i$  is the outcome of a numerical experiment or another type of experiment and is thus only an approximation to the true value. It is then always accompanied by an error estimate, often limited to a statistical error estimate given by the standard deviation discussed earlier. In the discussion here we will treat  $y_i$  as our exact value for the response variable.

Introducing the standard deviation  $\sigma_i$  for each measurement  $y_i$ , we define now the  $\chi^2$  function (omitting the  $1/n$  term) as

$$\chi^2(\boldsymbol{\beta}) = \frac{1}{n} \sum_{i=0}^{n-1} \frac{(y_i - \tilde{y}_i)^2}{\sigma_i^2} = \frac{1}{n} \left\{ (\mathbf{y} - \tilde{\mathbf{y}})^T \frac{1}{\boldsymbol{\Sigma}^2} (\mathbf{y} - \tilde{\mathbf{y}}) \right\},$$

where the matrix  $\boldsymbol{\Sigma}$  is a diagonal matrix with  $\sigma_i$  as matrix elements.

## The $\chi^2$ function

In order to find the parameters  $\beta_i$  we will then minimize the spread of  $\chi^2(\boldsymbol{\beta})$  by requiring

$$\frac{\partial \chi^2(\boldsymbol{\beta})}{\partial \beta_j} = \frac{\partial}{\partial \beta_j} \left[ \frac{1}{n} \sum_{i=0}^{n-1} \left( \frac{y_i - \beta_0 x_{i,0} - \beta_1 x_{i,1} - \beta_2 x_{i,2} - \cdots - \beta_{n-1} x_{i,n-1}}{\sigma_i} \right)^2 \right] = 0,$$

which results in

$$\frac{\partial \chi^2(\boldsymbol{\beta})}{\partial \beta_j} = -\frac{2}{n} \left[ \sum_{i=0}^{n-1} \frac{x_{ij}}{\sigma_i} \left( \frac{y_i - \beta_0 x_{i,0} - \beta_1 x_{i,1} - \beta_2 x_{i,2} - \cdots - \beta_{n-1} x_{i,n-1}}{\sigma_i} \right) \right] = 0,$$

or in a matrix-vector form as

$$\frac{\partial \chi^2(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = 0 = \mathbf{A}^T (\mathbf{b} - \mathbf{A}\boldsymbol{\beta}).$$

where we have defined the matrix  $\mathbf{A} = \mathbf{X}/\boldsymbol{\Sigma}$  with matrix elements  $a_{ij} = x_{ij}/\sigma_i$  and the vector  $\mathbf{b}$  with elements  $b_i = y_i/\sigma_i$ .

### The $\chi^2$ function

We can rewrite

$$\frac{\partial \chi^2(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = 0 = \mathbf{A}^T (\mathbf{b} - \mathbf{A}\boldsymbol{\beta}),$$

as

$$\mathbf{A}^T \mathbf{b} = \mathbf{A}^T \mathbf{A} \boldsymbol{\beta},$$

and if the matrix  $\mathbf{A}^T \mathbf{A}$  is invertible we have the solution

$$\boldsymbol{\beta} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}.$$

### The $\chi^2$ function

If we then introduce the matrix

$$\mathbf{H} = (\mathbf{A}^T \mathbf{A})^{-1},$$

we have then the following expression for the parameters  $\beta_j$  (the matrix elements of  $\mathbf{H}$  are  $h_{ij}$ )

$$\beta_j = \sum_{k=0}^{p-1} h_{jk} \sum_{i=0}^{n-1} \frac{y_i}{\sigma_i} \frac{x_{ik}}{\sigma_i} = \sum_{k=0}^{p-1} h_{jk} \sum_{i=0}^{n-1} b_i a_{ik}$$

We state without proof the expression for the uncertainty in the parameters  $\beta_j$  as (we leave this as an exercise)

$$\sigma^2(\beta_j) = \sum_{i=0}^{n-1} \sigma_i^2 \left( \frac{\partial \beta_j}{\partial y_i} \right)^2,$$

resulting in

$$\sigma^2(\beta_j) = \left( \sum_{k=0}^{p-1} h_{jk} \sum_{i=0}^{n-1} a_{ik} \right) \left( \sum_{l=0}^{p-1} h_{jl} \sum_{m=0}^{n-1} a_{ml} \right) = h_{jj}.$$

## The $\chi^2$ function

The first step here is to approximate the function  $y$  with a first-order polynomial, that is we write

$$y = y(x) \rightarrow y(x_i) \approx \beta_0 + \beta_1 x_i.$$

By computing the derivatives of  $\chi^2$  with respect to  $\beta_0$  and  $\beta_1$  show that these are given by

$$\frac{\partial \chi^2(\beta)}{\partial \beta_0} = -2 \left[ \frac{1}{n} \sum_{i=0}^{n-1} \left( \frac{y_i - \beta_0 - \beta_1 x_i}{\sigma_i^2} \right) \right] = 0,$$

and

$$\frac{\partial \chi^2(\beta)}{\partial \beta_1} = -\frac{2}{n} \left[ \sum_{i=0}^{n-1} x_i \left( \frac{y_i - \beta_0 - \beta_1 x_i}{\sigma_i^2} \right) \right] = 0.$$

## The $\chi^2$ function

For a linear fit (a first-order polynomial) we don't need to invert a matrix!! Defining

$$\gamma = \sum_{i=0}^{n-1} \frac{1}{\sigma_i^2},$$

$$\gamma_x = \sum_{i=0}^{n-1} \frac{x_i}{\sigma_i^2},$$

$$\gamma_y = \sum_{i=0}^{n-1} \left( \frac{y_i}{\sigma_i^2} \right),$$

$$\gamma_{xx} = \sum_{i=0}^{n-1} \frac{x_i x_i}{\sigma_i^2},$$

$$\gamma_{xy} = \sum_{i=0}^{n-1} \frac{y_i x_i}{\sigma_i^2},$$

we obtain

$$\beta_0 = \frac{\gamma_{xx} \gamma_y - \gamma_x \gamma_{xy}}{\gamma \gamma_{xx} - \gamma_x^2},$$

$$\beta_1 = \frac{\gamma_{xy} \gamma - \gamma_x \gamma_y}{\gamma \gamma_{xx} - \gamma_x^2}.$$

This approach (different linear and non-linear regression) suffers often from both being underdetermined and overdetermined in the unknown coefficients  $\beta_i$ . A better approach is to use the Singular Value Decomposition (SVD) method discussed below. Or using Lasso and Ridge regression. See below.

## Regression Examples

### Fitting an Equation of State for Dense Nuclear Matter

Before we continue, let us introduce yet another example. We are going to fit the nuclear equation of state using results from many-body calculations. The equation of state we have made available here, as function of density, has been derived using modern nucleon-nucleon potentials with [the addition of three-body interactions](#). This time the file is presented as a standard `csv` file.

The beginning of the Python code here is similar to what you have seen before, with the same initializations and declarations. We use also `pandas` again, rather extensively in order to organize our data.

The difference now is that we use **Scikit-Learn's** regression tools instead of our own matrix inversion implementation. Furthermore, we sneak in **Ridge** regression (to be discussed below) which includes a hyperparameter  $\lambda$ , also to be explained below.

### The code

The above simple polynomial in density  $\rho$  gives an excellent fit to the data.

We note also that there is a small deviation between the standard OLS and the Ridge regression at higher densities. We discuss this in more detail below.

### Splitting our Data in Training and Test data

It is normal in essentially all Machine Learning studies to split the data in a training set and a test set (sometimes also an additional validation set). **Scikit-Learn** has an own function for this. There is no explicit recipe for how much data should be included as training data and say test data. An accepted rule of thumb is to use approximately 2/3 to 4/5 of the data as training data. We will postpone a discussion of this splitting to the end of these notes and our discussion of the so-called **bias-variance** tradeoff. Here we limit ourselves to repeat the above equation of state fitting example but now splitting the data into a training set and a test set.

### The Boston housing data example

The Boston housing data set was originally a part of UCI Machine Learning Repository and has been removed now. The data set is now included in **Scikit-Learn's** library. There are 506 samples and 13 feature (predictor) variables in this data set. The objective is to predict the value of prices of the house using the features (predictors) listed here.

The features/predictors are

1. CRIM: Per capita crime rate by town
2. ZN: Proportion of residential land zoned for lots over 25000 square feet

3. INDUS: Proportion of non-retail business acres per town
4. CHAS: Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
5. NOX: Nitric oxide concentration (parts per 10 million)
6. RM: Average number of rooms per dwelling
7. AGE: Proportion of owner-occupied units built prior to 1940
8. DIS: Weighted distances to five Boston employment centers
9. RAD: Index of accessibility to radial highways
10. TAX: Full-value property tax rate per USD10000
11. B:  $1000(Bk - 0.63)^2$ , where  $Bk$  is the proportion of [people of African American descent] by town
12. LSTAT: Percentage of lower status of the population
13. MEDV: Median value of owner-occupied homes in USD 1000s

## Housing data, the code

We start by importing the libraries and load the Boston Housing DataSet from **Scikit-Learn**

Then we invoke Pandas and preprocess the data We can then visualize the data

It is now useful to look at the correlation matrix From the above coorelation plot we can see that **MEDV** is strongly correlated to **LSTAT** and **RM**. We see also that **RAD** and **TAX** are stronly correlated, but we don't include this in our features together to avoid multi-collinearity

Now we start training our model We split the data into training and test sets Then we use the linear regression functionality from **Scikit-Learn**

## Singular Value Decomposition Algorithm

### The singular value decomposition

The examples we have looked at so far are cases where we normally can invert the matrix  $\mathbf{X}^T \mathbf{X}$ . Using a polynomial expansion as we did both for the masses and the fitting of the equation of state, leads to row vectors of the design matrix which are essentially orthogonal due to the polynomial character of our model. Obtaining the inverse of the design matrix is then often done via a so-called LU, QR or Cholesky decomposition.

This may however not be the case in general and a standard matrix inversion algorithm based on say LU, QR or Cholesky decomposition may lead to singularities. We will see examples of this below.

There is however a way to partially circumvent this problem and also gain some insight about the ordinary least squares approach.

This is given by the **Singular Value Decomposition** algorithm, perhaps the most powerful linear algebra algorithm. Let us look at a different example where we may have problems with the standard matrix inversion algorithm. Thereafter we dive into the math of the SVD.

## Linear Regression Problems

One of the typical problems we encounter with linear regression, in particular when the matrix  $\mathbf{X}$  (our so-called design matrix) is high-dimensional, are problems with near singular or singular matrices. The column vectors of  $\mathbf{X}$  may be linearly dependent, normally referred to as super-collinearity. This means that the matrix may be rank deficient and it is basically impossible to model the data using linear regression. As an example, consider the matrix

$$\mathbf{X} = \begin{bmatrix} 1 & -1 & 2 \\ 1 & 0 & 1 \\ 1 & 2 & -1 \\ 1 & 1 & 0 \end{bmatrix}$$

The columns of  $\mathbf{X}$  are linearly dependent. We see this easily since the first column is the row-wise sum of the other two columns. The rank (more correct, the column rank) of a matrix is the dimension of the space spanned by the column vectors. Hence, the rank of  $\mathbf{X}$  is equal to the number of linearly independent columns. In this particular case the matrix has rank 2.

Super-collinearity of an  $(n \times p)$ -dimensional design matrix  $\mathbf{X}$  implies that the inverse of the matrix  $\mathbf{X}^T \mathbf{X}$  (the matrix we need to invert to solve the linear regression equations) is non-invertible. If we have a square matrix that does not have an inverse, we say this matrix singular. The example here demonstrates this

$$\mathbf{X} = \begin{bmatrix} 1 & -1 \\ 1 & -1 \end{bmatrix}.$$

We see easily that  $\det(\mathbf{X}) = x_{11}x_{22} - x_{12}x_{21} = 1 \times (-1) - 1 \times (-1) = 0$ . Hence,  $\mathbf{X}$  is singular and its inverse is undefined. This is equivalent to saying that the matrix  $\mathbf{X}$  has at least an eigenvalue which is zero.

## Fixing the singularity

If our design matrix  $\mathbf{X}$  which enters the linear regression problem

$$\boldsymbol{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}, \quad (1)$$

has linearly dependent column vectors, we will not be able to compute the inverse of  $\mathbf{X}^T \mathbf{X}$  and we cannot find the parameters (estimators)  $\beta_i$ . The estimators are only well-defined if  $(\mathbf{X}^T \mathbf{X})^{-1}$  exists. This is more likely to happen when the

matrix  $\mathbf{X}$  is high-dimensional. In this case it is likely to encounter a situation where the regression parameters  $\beta_i$  cannot be estimated.

A cheap *ad hoc* approach is simply to add a small diagonal component to the matrix to invert, that is we change

$$\mathbf{X}^T \mathbf{X} \rightarrow \mathbf{X}^T \mathbf{X} + \lambda \mathbf{I},$$

where  $\mathbf{I}$  is the identity matrix. When we discuss **Ridge** regression this is actually what we end up evaluating. The parameter  $\lambda$  is called a hyperparameter. More about this later.

## Basic math of the SVD

From standard linear algebra we know that a square matrix  $\mathbf{X}$  can be diagonalized if and only if it is a so-called **normal matrix**, that is if  $\mathbf{X} \in \mathbb{R}^{n \times n}$  we have  $\mathbf{X}\mathbf{X}^T = \mathbf{X}^T\mathbf{X}$  or if  $\mathbf{X} \in \mathbb{C}^{n \times n}$  we have  $\mathbf{X}\mathbf{X}^\dagger = \mathbf{X}^\dagger\mathbf{X}$ . The matrix has then a set of eigenpairs

$(\lambda_1, \mathbf{u}_1), \dots, (\lambda_n, \mathbf{u}_n)$ , and the eigenvalues are given by the diagonal matrix  $\mathbf{\Sigma} = \text{Diag}(\lambda_1, \dots, \lambda_n)$ .

The matrix  $\mathbf{X}$  can be written in terms of an orthogonal/unitary transformation  $\mathbf{U}$

$$\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T,$$

with  $\mathbf{U}\mathbf{U}^T = \mathbf{I}$  or  $\mathbf{U}\mathbf{U}^\dagger = \mathbf{I}$ .

Not all square matrices are diagonalizable. A matrix like the one discussed above

$$\mathbf{X} = \begin{bmatrix} 1 & -1 \\ 1 & -1 \end{bmatrix}$$

is not diagonalizable, it is a so-called **defective matrix**. It is easy to see that the condition  $\mathbf{X}\mathbf{X}^T = \mathbf{X}^T\mathbf{X}$  is not fulfilled.

## The SVD, a Fantastic Algorithm

However, and this is the strength of the SVD algorithm, any general matrix  $\mathbf{X}$  can be decomposed in terms of a diagonal matrix and two orthogonal/unitary matrices. The **Singular Value Decomposition (SVD) theorem** states that a general  $m \times n$  matrix  $\mathbf{X}$  can be written in terms of a diagonal matrix  $\mathbf{\Sigma}$  of dimensionality  $n \times n$  and two orthogonal matrices  $\mathbf{U}$  and  $\mathbf{V}$ , where the first has dimensionality  $m \times m$  and the last dimensionality  $n \times n$ . We have then

$$\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

As an example, the above defective matrix can be decomposed as

$$\mathbf{X} = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix} \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T,$$

with eigenvalues  $\sigma_1 = 2$  and  $\sigma_2 = 0$ . The SVD exists always!

## Another Example

Consider the following matrix which can be SVD decomposed as

$$\mathbf{X} = \frac{1}{15} \begin{bmatrix} 14 & 2 \\ 4 & 22 \\ 16 & 13 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 & 2 & 2 \\ 2 & -1 & 1 \\ 2 & 1 & -2 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 1 \\ 0 & 0 \end{bmatrix} \frac{1}{5} \begin{bmatrix} 3 & 4 \\ 4 & -3 \end{bmatrix} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T.$$

This is a  $3 \times 2$  matrix which is decomposed in terms of a  $3 \times 3$  matrix  $\mathbf{U}$ , and a  $2 \times 2$  matrix  $\mathbf{V}$ . It is easy to see that  $\mathbf{U}$  and  $\mathbf{V}$  are orthogonal (how?).

And the SVD decomposition (singular values) gives eigenvalues  $\sigma_i \geq \sigma_{i+1}$  for all  $i$  and for dimensions larger than  $i = 2$ , the eigenvalues (singular values) are zero.

In the general case, where our design matrix  $\mathbf{X}$  has dimension  $n \times p$ , the matrix is thus decomposed into an  $n \times n$  orthogonal matrix  $\mathbf{U}$ , a  $p \times p$  orthogonal matrix  $\mathbf{V}$  and a diagonal matrix  $\mathbf{\Sigma}$  with  $r = \min(n, p)$  singular values  $\sigma_i \geq 0$  on the main diagonal and zeros filling the rest of the matrix. There are at most  $p$  singular values assuming that  $n > p$ . In our regression examples for the nuclear masses and the equation of state this is indeed the case, while for the Ising model we have  $p > n$ . These are often cases that lead to near singular or singular matrices.

The columns of  $\mathbf{U}$  are called the left singular vectors while the columns of  $\mathbf{V}$  are the right singular vectors.

## Economy-size SVD

If we assume that  $n > p$ , then our matrix  $\mathbf{U}$  has dimension  $n \times n$ . The last  $n - p$  columns of  $\mathbf{U}$  become however irrelevant in our calculations since they are multiplied with the zeros in  $\mathbf{\Sigma}$ .

The economy-size decomposition removes extra rows or columns of zeros from the diagonal matrix of singular values,  $\mathbf{\Sigma}$ , along with the columns in either  $\mathbf{U}$  or  $\mathbf{V}$  that multiply those zeros in the expression. Removing these zeros and columns can improve execution time and reduce storage requirements without compromising the accuracy of the decomposition.

If  $n > p$ , we keep only the first  $p$  columns of  $\mathbf{U}$  and  $\mathbf{\Sigma}$  has dimension  $p \times p$ . If  $p > n$ , then only the first  $n$  columns of  $\mathbf{V}$  are computed and  $\mathbf{\Sigma}$  has dimension  $n \times n$ . The  $n = p$  case is obvious, we retain the full SVD. In general the economy-size SVD leads to less FLOPS and still conserving the desired accuracy.

## Codes for the SVD

The matrix  $\mathbf{X}$  has columns that are linearly dependent. The first column is the row-wise sum of the other two columns. The rank of a matrix (the column rank) is the dimension of space spanned by the column vectors. The rank of the matrix is the number of linearly independent columns, in this case just 2. We see this from the singular values when running the above code. Running



the standard inversion algorithm for matrix inversion with  $\mathbf{X}^T \mathbf{X}$  results in the program terminating due to a singular matrix.

## Mathematical Properties

There are several interesting mathematical properties which will be relevant when we are going to discuss the differences between say ordinary least squares (OLS) and **Ridge** regression.

We have from OLS that the parameters of the linear approximation are given by

$$\tilde{\mathbf{y}} = \mathbf{X}\boldsymbol{\beta} = \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

The matrix to invert can be rewritten in terms of our SVD decomposition as

$$\mathbf{X}^T \mathbf{X} = \mathbf{V} \boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T.$$

Using the orthogonality properties of  $\mathbf{U}$  we have

$$\mathbf{X}^T \mathbf{X} = \mathbf{V} \boldsymbol{\Sigma}^T \boldsymbol{\Sigma} \mathbf{V}^T = \mathbf{V} \mathbf{D} \mathbf{V}^T,$$

with  $\mathbf{D}$  being a diagonal matrix with values along the diagonal given by the singular values squared.

This means that

$$(\mathbf{X}^T \mathbf{X}) \mathbf{V} = \mathbf{V} \mathbf{D},$$

that is the eigenvectors of  $(\mathbf{X}^T \mathbf{X})$  are given by the columns of the right singular matrix of  $\mathbf{X}$  and the eigenvalues are the squared singular values. It is easy to show (show this) that

$$(\mathbf{X} \mathbf{X}^T) \mathbf{U} = \mathbf{U} \mathbf{D},$$

that is, the eigenvectors of  $(\mathbf{X} \mathbf{X}^T)$  are the columns of the left singular matrix and the eigenvalues are the same.

Going back to our OLS equation we have

$$\mathbf{X}\boldsymbol{\beta} = \mathbf{X} (\mathbf{V} \mathbf{D} \mathbf{V}^T)^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T (\mathbf{V} \mathbf{D} \mathbf{V}^T)^{-1} (\mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T)^T \mathbf{y} = \mathbf{U} \mathbf{U}^T \mathbf{y}.$$

We will come back to this expression when we discuss Ridge regression.

## Beyond Ordinary Least Squares

### Ridge and LASSO Regression

Let us remind ourselves about the expression for the standard Mean Squared Error (MSE) which we used to define our cost function and the equations for the ordinary least squares (OLS) method, that is our optimization problem is

$$\min_{\boldsymbol{\beta} \in \mathbb{R}^p} \frac{1}{n} \left\{ (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right\}.$$

or we can state it as

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{n} \sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2 = \frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2,$$

where we have used the definition of a norm-2 vector, that is

$$\|\mathbf{x}\|_2 = \sqrt{\sum_i x_i^2}.$$

By minimizing the above equation with respect to the parameters  $\beta$  we could then obtain an analytical expression for the parameters  $\beta$ . We can add a regularization parameter  $\lambda$  by defining a new cost function to be optimized, that is

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2$$

which leads to the Ridge regression minimization problem where we require that  $\|\beta\|_2^2 \leq t$ , where  $t$  is a finite number larger than zero. By defining

$$C(\mathbf{X}, \beta) = \frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1,$$

we have a new optimization equation

$$\min_{\beta \in \mathbb{R}^p} \frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1$$

which leads to Lasso regression. Lasso stands for least absolute shrinkage and selection operator.

Here we have defined the norm-1 as

$$\|\mathbf{x}\|_1 = \sum_i |x_i|.$$

## More on Ridge Regression

Using the matrix-vector expression for Ridge regression (we drop the  $1/n$  factor),

$$C(\mathbf{X}, \beta) = \{(\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)\} + \lambda \beta^T \beta,$$

by taking the derivatives with respect to  $\beta$  we obtain

$$\frac{\partial C(\beta)}{\partial \beta} = 0 = 2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta) - 2\lambda \beta.$$

We obtain a slightly modified matrix inversion problem which for finite values of  $\lambda$  does not suffer from singularity problems, that is

$$\beta^{\text{Ridge}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y},$$

with  $\mathbf{I}$  being a  $p \times p$  identity matrix with the constraint that

$$\sum_{i=0}^{p-1} \beta_i^2 \leq t,$$

with  $t$  a finite positive number.

We see that Ridge regression is nothing but the standard OLS with a modified diagonal term added to  $\mathbf{X}^T \mathbf{X}$ . The consequences, in particular for our discussion of the bias-variance tradeoff are rather interesting.

Furthermore, if we use the result above in terms of the SVD decomposition (our analysis was done for the OLS method), we get

$$(\mathbf{X} \mathbf{X}^T) \mathbf{U} = \mathbf{U} \mathbf{D}.$$

We can analyse the OLS solutions in terms of the eigenvectors (the columns) of the right singular value matrix  $\mathbf{U}$  as

$$\mathbf{X} \boldsymbol{\beta} = \mathbf{X} (\mathbf{V} \mathbf{D} \mathbf{V}^T)^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T (\mathbf{V} \mathbf{D} \mathbf{V}^T)^{-1} (\mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T)^T \mathbf{y} = \mathbf{U} \mathbf{U}^T \mathbf{y}$$

For Ridge regression this becomes

$$\mathbf{X} \boldsymbol{\beta}^{\text{Ridge}} = \mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T (\mathbf{V} \mathbf{D} \mathbf{V}^T + \lambda \mathbf{I})^{-1} (\mathbf{U} \boldsymbol{\Sigma} \mathbf{V}^T)^T \mathbf{y} = \sum_{j=0}^{p-1} \mathbf{u}_j \mathbf{u}_j^T \frac{\sigma_j^2}{\sigma_j^2 + \lambda} \mathbf{y},$$

with the vectors  $\mathbf{u}_j$  being the columns of  $\mathbf{U}$ .

## Interpreting the Ridge results

Since  $\lambda \geq 0$ , it means that compared to OLS, we have

$$\frac{\sigma_j^2}{\sigma_j^2 + \lambda} \leq 1.$$

Ridge regression finds the coordinates of  $\mathbf{y}$  with respect to the orthonormal basis  $\mathbf{U}$ , it then shrinks the coordinates by  $\frac{\sigma_j^2}{\sigma_j^2 + \lambda}$ . Recall that the SVD has eigenvalues ordered in a descending way, that is  $\sigma_i \geq \sigma_{i+1}$ .

For small eigenvalues  $\sigma_i$  it means that their contributions become less important, a fact which can be used to reduce the number of degrees of freedom. Actually, calculating the variance of  $\mathbf{X} \mathbf{v}_j$  shows that this quantity is equal to  $\sigma_j^2/n$ . With a parameter  $\lambda$  we can thus shrink the role of specific parameters.

## More interpretations

For the sake of simplicity, let us assume that the design matrix is orthonormal, that is

$$\mathbf{X}^T \mathbf{X} = (\mathbf{X}^T \mathbf{X})^{-1} = \mathbf{I}.$$

In this case the standard OLS results in

$$\boldsymbol{\beta}^{\text{OLS}} = \mathbf{X}^T \mathbf{y} = \sum_{i=0}^{p-1} \mathbf{u}_i \mathbf{u}_i^T \mathbf{y},$$

and

$$\boldsymbol{\beta}^{\text{Ridge}} = (\mathbf{I} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y} = (1 + \lambda)^{-1} \boldsymbol{\beta}^{\text{OLS}},$$

that is the Ridge estimator scales the OLS estimator by the inverse of a factor  $1 + \lambda$ , and the Ridge estimator converges to zero when the hyperparameter goes to infinity.

For more discussions of Ridge and Lasso regression, [Wessel van Wieringen's](#) article is highly recommended. Similarly, [Mehta et al's article](#) is also recommended.

## Statistics and Resampling Techniques

### Where are we going?

Before we proceed, we need to rethink what we have been doing. In our eager to fit the data, we have omitted several important elements in our regression analysis. In what follows we will

1. remind ourselves about some statistical properties, including a discussion of mean values, variance and the so-called bias-variance tradeoff
2. introduce resampling techniques like cross-validation, bootstrapping and jackknife and more

This will allow us to link the standard linear algebra methods we have discussed above to a statistical interpretation of the methods.

### Resampling methods

Resampling methods are an indispensable tool in modern statistics. They involve repeatedly drawing samples from a training set and refitting a model of interest on each sample in order to obtain additional information about the fitted model. For example, in order to estimate the variability of a linear regression fit, we can repeatedly draw different samples from the training data, fit a linear regression to each new sample, and then examine the extent to which the resulting fits differ. Such an approach may allow us to obtain information that would not be available from fitting the model only once using the original training sample.

Two resampling methods are often used in Machine Learning analyses,

1. The **bootstrap method**
2. and **Cross-Validation**

In addition there are several other methods such as the Jackknife and the Blocking methods. We will discuss in particular cross-validation and the bootstrap method.

## Resampling approaches can be computationally expensive

Resampling approaches can be computationally expensive, because they involve fitting the same statistical method multiple times using different subsets of the training data. However, due to recent advances in computing power, the computational requirements of resampling methods generally are not prohibitive. In this chapter, we discuss two of the most commonly used resampling methods, cross-validation and the bootstrap. Both methods are important tools in the practical application of many statistical learning procedures. For example, cross-validation can be used to estimate the test error associated with a given statistical learning method in order to evaluate its performance, or to select the appropriate level of flexibility. The process of evaluating a model's performance is known as model assessment, whereas the process of selecting the proper level of flexibility for a model is known as model selection. The bootstrap is widely used.

## Why resampling methods ?

Statistical analysis.

- Our simulations can be treated as *computer experiments*. This is particularly the case for Monte Carlo methods
- The results can be analysed with the same statistical tools as we would use analysing experimental data.
- As in all experiments, we are looking for expectation values and an estimate of how accurate they are, i.e., possible sources for errors.

## Statistical analysis

- As in other experiments, many numerical experiments have two classes of errors:
  - Statistical errors
  - Systematical errors
- Statistical errors can be estimated using standard tools from statistics
- Systematical errors are method specific and must be treated differently from case to case.

## Linking the regression analysis with a statistical interpretation

We are going to discuss several statistical properties which can be obtained in terms of analytical expressions. The advantage of doing linear regression is that we actually end up with analytical expressions for several statistical quantities. Standard least squares and Ridge regression allow us to derive quantities like the variance and other expectation values in a rather straightforward way.

It is assumed that  $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$  and the  $\varepsilon_i$  are independent, i.e.:

$$\text{Cov}(\varepsilon_{i_1}, \varepsilon_{i_2}) = \begin{cases} \sigma^2 & \text{if } i_1 = i_2, \\ 0 & \text{if } i_1 \neq i_2. \end{cases}$$

The randomness of  $\varepsilon_i$  implies that  $\mathbf{y}_i$  is also a random variable. In particular,  $\mathbf{y}_i$  is normally distributed, because  $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$  and  $\mathbf{X}_{i,*}\boldsymbol{\beta}$  is a non-random scalar. To specify the parameters of the distribution of  $\mathbf{y}_i$  we need to calculate its first two moments.

Recall that  $\mathbf{X}$  is a matrix of dimensionality  $n \times p$ . The notation above  $\mathbf{X}_{i,*}$  means that we are looking at the row number  $i$  and perform a sum over all values  $p$ .

## Assumptions made

The assumption we have made here can be summarized as (and this is going to be useful when we discuss the bias-variance trade off) that there exists a function  $f(\mathbf{x})$  and a normal distributed error  $\varepsilon \sim \mathcal{N}(0, \sigma^2)$  which describe our data

$$\mathbf{y} = f(\mathbf{x}) + \varepsilon$$

We approximate this function with our model from the solution of the linear regression equations, that is our function  $f$  is approximated by  $\tilde{\mathbf{y}}$  where we want to minimize  $(\mathbf{y} - \tilde{\mathbf{y}})^2$ , our MSE, with

$$\tilde{\mathbf{y}} = \mathbf{X}\boldsymbol{\beta} \approx f(\mathbf{x}).$$

Note that we reserve the design matrix  $\mathbf{X}$  to represent our specific rewrite of the input variables  $\mathbf{x}$ .

## Expectation value and variance

We can calculate the expectation value of  $\mathbf{y}$  for a given element  $i$

$$\mathbb{E}(y_i) = \mathbb{E}(\mathbf{X}_{i,*}\boldsymbol{\beta}) + \mathbb{E}(\varepsilon_i) = \mathbf{X}_{i,*}\boldsymbol{\beta},$$

while its variance is

$$\begin{aligned}
\text{Var}(y_i) &= \mathbb{E}\{[y_i - \mathbb{E}(y_i)]^2\} = \mathbb{E}(y_i^2) - [\mathbb{E}(y_i)]^2 \\
&= \mathbb{E}[(\mathbf{X}_{i,*} \boldsymbol{\beta} + \varepsilon_i)^2] - (\mathbf{X}_{i,*} \boldsymbol{\beta})^2 \\
&= \mathbb{E}[(\mathbf{X}_{i,*} \boldsymbol{\beta})^2 + 2\varepsilon_i \mathbf{X}_{i,*} \boldsymbol{\beta} + \varepsilon_i^2] - (\mathbf{X}_{i,*} \boldsymbol{\beta})^2 \\
&= (\mathbf{X}_{i,*} \boldsymbol{\beta})^2 + 2\mathbb{E}(\varepsilon_i) \mathbf{X}_{i,*} \boldsymbol{\beta} + \mathbb{E}(\varepsilon_i^2) - (\mathbf{X}_{i,*} \boldsymbol{\beta})^2 \\
&= \mathbb{E}(\varepsilon_i^2) = \text{Var}(\varepsilon_i) = \sigma^2.
\end{aligned}$$

Hence,  $y_i \sim \mathcal{N}(\mathbf{X}_{i,*} \boldsymbol{\beta}, \sigma^2)$ , that is  $\mathbf{y}$  follows a normal distribution with mean value  $\mathbf{X}\boldsymbol{\beta}$  and variance  $\sigma^2$  (not be confused with the singular values of the SVD).

### Expectation value and variance for $\boldsymbol{\beta}$

With the OLS expressions for the parameters  $\boldsymbol{\beta}$  we can evaluate the expectation value

$$\mathbb{E}(\boldsymbol{\beta}) = \mathbb{E}[(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}] = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbb{E}[\mathbf{Y}] = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} \boldsymbol{\beta} = \boldsymbol{\beta}.$$

This means that the estimator of the regression parameters is unbiased.

We can also calculate the variance

The variance of  $\boldsymbol{\beta}$  is

$$\begin{aligned}
\text{Var}(\boldsymbol{\beta}) &= \mathbb{E}\{[\boldsymbol{\beta} - \mathbb{E}(\boldsymbol{\beta})][\boldsymbol{\beta} - \mathbb{E}(\boldsymbol{\beta})]^\top\} \\
&= \mathbb{E}\{[(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y} - \boldsymbol{\beta}][(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y} - \boldsymbol{\beta}]^\top\} \\
&= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbb{E}\{\mathbf{Y} \mathbf{Y}^\top\} \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} - \boldsymbol{\beta} \boldsymbol{\beta}^\top \\
&= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \{\mathbf{X} \boldsymbol{\beta} \boldsymbol{\beta}^\top \mathbf{X}^\top + \sigma^2 \mathbf{I}_n\} \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1} - \boldsymbol{\beta} \boldsymbol{\beta}^\top \\
&= \boldsymbol{\beta} \boldsymbol{\beta}^\top + \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1} - \boldsymbol{\beta} \boldsymbol{\beta}^\top = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1},
\end{aligned}$$

where we have used that  $\mathbb{E}(\mathbf{Y} \mathbf{Y}^\top) = \mathbf{X} \boldsymbol{\beta} \boldsymbol{\beta}^\top \mathbf{X}^\top + \sigma^2 \mathbf{I}_n$ . From  $\text{Var}(\boldsymbol{\beta}) = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1}$ , one obtains an estimate of the variance of the estimate of the  $j$ -th regression coefficient:  $\sigma^2(\hat{\beta}_j) = \sigma^2 \sqrt{[(\mathbf{X}^\top \mathbf{X})^{-1}]_{jj}}$ . This may be used to construct a confidence interval for the estimates.

In a similar way, we can obtain analytical expressions for say the expectation values of the parameters  $\boldsymbol{\beta}$  and their variance when we employ Ridge regression, allowing us again to define a confidence interval.

It is rather straightforward to show that

$$\mathbb{E}[\boldsymbol{\beta}^{\text{Ridge}}] = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{pp})^{-1} (\mathbf{X}^\top \mathbf{X}) \boldsymbol{\beta}^{\text{OLS}}.$$

We see clearly that  $\mathbb{E}[\boldsymbol{\beta}^{\text{Ridge}}] \neq \boldsymbol{\beta}^{\text{OLS}}$  for any  $\lambda > 0$ . We say then that the ridge estimator is biased.

We can also compute the variance as

$$\text{Var}[\boldsymbol{\beta}^{\text{Ridge}}] = \sigma^2 [\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}]^{-1} \mathbf{X}^\top \mathbf{X} \{[\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}]^{-1}\}^\top,$$

and it is easy to see that if the parameter  $\lambda$  goes to infinity then the variance of Ridge parameters  $\beta$  goes to zero.

With this, we can compute the difference

$$\text{Var}[\beta^{\text{OLS}}] - \text{Var}(\beta^{\text{Ridge}}) = \sigma^2 [\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}]^{-1} [2\lambda \mathbf{I} + \lambda^2 (\mathbf{X}^T \mathbf{X})^{-1}] \{[\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}]^{-1}\}^T.$$

The difference is non-negative definite since each component of the matrix product is non-negative definite. This means the variance we obtain with the standard OLS will always for  $\lambda > 0$  be larger than the variance of  $\beta$  obtained with the Ridge estimator. This has interesting consequences when we discuss the so-called bias-variance trade-off below.

## Resampling methods

### Resampling methods, basic overview

With all these analytical equations for both the OLS and Ridge regression, we will now outline how to assess a given model. This will lead us to a discussion of the so-called bias-variance tradeoff (see below) and so-called resampling methods.

One of the quantities we have discussed as a way to measure errors is the mean-squared error (MSE), mainly used for fitting of continuous functions. Another choice is the absolute error.

In the discussions below we will focus on the MSE and in particular since we will split the data into test and training data, we discuss the

1. prediction error or simply the **test error**, where we have a fixed training set and the test error is the MSE arising from the data reserved for testing. We discuss also the
2. training error  $\text{Err}_{\text{Train}}$ , which is the average loss over the training data.

As our model becomes more and more complex, more of the training data tends to be used. The training may then adapt to more complicated structures in the data. This may lead to a decrease in the bias (see below for code example) and a slight increase of the variance for the test error. For a certain level of complexity the test error will reach a minimum, before starting to increase again. The training error reaches a saturation.

### Train and test, an example

### Bringing back our Equation of State data

### Resampling methods: Jackknife and Bootstrap

Two famous resampling methods are the **independent bootstrap** and the **jackknife**.

The jackknife is a special case of the independent bootstrap. Still, the jackknife was made popular prior to the independent bootstrap. And as the



popularity of the independent bootstrap soared, new variants, such as **the dependent bootstrap**.

The Jackknife and independent bootstrap work for independent, identically distributed random variables. If these conditions are not satisfied, the methods will fail.

## Resampling methods: Bootstrap

Bootstrapping is a nonparametric approach to statistical inference that substitutes computation for more traditional distributional assumptions and asymptotic results. Bootstrapping offers a number of advantages:

1. The bootstrap is quite general, although there are some cases in which it fails.
2. Because it does not require distributional assumptions (such as normally distributed errors), the bootstrap can provide more accurate inferences when the data are not well behaved or when the sample size is small.
3. It is possible to apply the bootstrap to statistics with sampling distributions that are difficult to derive, even asymptotically.
4. It is relatively simple to apply the bootstrap.

## Resampling methods: Bootstrap steps

The independent bootstrap works like this:

1. Draw with replacement  $n$  numbers for the observed variables  $\mathbf{x} = (x_1, x_2, \dots, x_n)$ .
2. Define a vector  $\mathbf{x}^*$  containing the values which were drawn from  $\mathbf{x}$ .
3. Using the vector  $\mathbf{x}^*$  compute the estimate (parameter)  $\hat{\theta}^*$  by evaluating  $\hat{\theta}$  under the observations  $\mathbf{x}^*$ .
4. Repeat this process  $k$  times.

When you are done, you can draw a histogram of the relative frequency of the estimator/parameter  $\hat{\theta}^*$ . This is your estimate of the probability distribution  $p(x)$ . Using this probability distribution you can estimate any statistics thereof. In principle you never draw the histogram of the relative frequency of  $\hat{\theta}^*$ . Instead you use the estimators corresponding to the statistic of interest. For example, if you are interested in estimating the variance of  $\hat{\theta}$ , apply the estimator  $\hat{\sigma}^2$  to the values  $\hat{\theta}^*$ .

## Code example for the Bootstrap method

The following code starts with a Gaussian distribution with mean value  $\mu = 100$  and variance  $\sigma = 15$ . We use this to generate the data used in the bootstrap analysis. The bootstrap analysis returns a data set after a given number of bootstrap operations (as many as we have data points). This data set consists of estimated mean values for each bootstrap operation. The histogram generated by the bootstrap method shows that the distribution for these mean values is also a Gaussian, centered around the mean value  $\mu = 100$  but with standard deviation  $\sigma/\sqrt{n}$ , where  $n$  is the number of bootstrap samples (in this case the same as the number of original data points). The value of the standard deviation is what we expect from the central limit theorem.

## Various steps in cross-validation

When the repetitive splitting of the data set is done randomly, samples may accidentally end up in a fast majority of the splits in either training or test set. Such samples may have an unbalanced influence on either model building or prediction evaluation. To avoid this  $k$ -fold cross-validation structures the data splitting. The samples are divided into  $k$  more or less equally sized exhaustive and mutually exclusive subsets. In turn (at each split) one of these subsets plays the role of the test set while the union of the remaining subsets constitutes the training set. Such a splitting warrants a balanced representation of each sample in both training and test set over the splits. Still the division into the  $k$  subsets involves a degree of randomness. This may be fully excluded when choosing  $k = n$ . This particular case is referred to as leave-one-out cross-validation (LOOCV).

## How to set up the cross-validation for Ridge and/or Lasso

- Define a range of interest for the penalty parameter.
- Divide the data set into training and test set comprising samples  $\{1, \dots, n\} \setminus i$  and  $\{i\}$ , respectively.
- Fit the linear regression model by means of OLS, or Ridge or Lasso estimations for each  $\lambda$  in the grid using the training set, and the corresponding estimate of the error variance  $\sigma_{-i}^2(\lambda)$ , as

$$\beta_{-i}(\lambda) = (\mathbf{X}_{-i,*}^T \mathbf{X}_{-i,*} + \lambda \mathbf{I}_{pp})^{-1} \mathbf{X}_{-i,*}^T \mathbf{y}_{-i}$$

- Evaluate the prediction performance of these models on the test set by  $\log\{L[y_i, \mathbf{X}_{i,*}; \beta_{-i}(\lambda), \sigma_{-i}^2(\lambda)]\}$ . Or, by the prediction error  $|y_i - \mathbf{X}_{i,*} \beta_{-i}(\lambda)|$ , the relative error, the error squared or the R2 score function.
- Repeat the first three steps such that each sample plays the role of the test set once.

- Average the prediction performances of the test sets at each grid point of the penalty bias/parameter. It is an estimate of the prediction performance of the model corresponding to this value of the penalty parameter on novel data. It is defined as

$$\frac{1}{n} \sum_{i=1}^n \log\{L[y_i, \mathbf{X}_{i,*}; \beta_{-i}(\lambda), \sigma_{-i}^2(\lambda)]\}.$$

## Cross-validation in brief

For the various values of  $k$

1. shuffle the dataset randomly.
2. Split the dataset into  $k$  groups.
3. For each unique group:
  - (a) Decide which group to use as set for test data
  - (b) Take the remaining groups as a training data set
  - (c) Fit a model on the training set and evaluate it on the test set
  - (d) Retain the evaluation score and discard the model
4. Summarize the model using the sample of model evaluation scores

## Code Example for Cross-validation and $k$ -fold Cross-validation

The code here uses Ridge regression with cross-validation (CV) resampling and  $k$ -fold CV in order to fit a specific polynomial.

## How to decide upon the best model?

### The bias-variance tradeoff

We will discuss the bias-variance tradeoff in the context of continuous predictions such as regression. However, many of the intuitions and ideas discussed here also carry over to classification tasks. Consider a dataset  $\mathcal{L}$  consisting of the data  $\mathbf{X}_{\mathcal{L}} = \{(y_j, \mathbf{x}_j), j = 0 \dots n-1\}$ .

Let us assume that the true data is generated from a noisy model

$$\mathbf{y} = f(\mathbf{x}) + \epsilon$$

where  $\epsilon$  is normally distributed with mean zero and standard deviation  $\sigma^2$ .

In our derivation of the ordinary least squares method we defined then an approximation to the function  $f$  in terms of the parameters  $\beta$  and the design matrix  $\mathbf{X}$  which embody our model, that is  $\tilde{\mathbf{y}} = \mathbf{X}\beta$ .

Thereafter we found the parameters  $\beta$  by optimizing the mean-squared error via the so-called cost function

$$C(\mathbf{X}, \beta) = \frac{1}{n} \sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2 = \mathbb{E}[(\mathbf{y} - \tilde{\mathbf{y}})^2].$$

We can rewrite this as

$$\mathbb{E}[(\mathbf{y} - \tilde{\mathbf{y}})^2] = \frac{1}{n} \sum_i (f_i - \mathbb{E}[\tilde{\mathbf{y}}])^2 + \frac{1}{n} \sum_i (\tilde{y}_i - \mathbb{E}[\tilde{\mathbf{y}}])^2 + \sigma^2.$$

The three terms represent the square of the bias of the learning method, which can be thought of as the error caused by the simplifying assumptions built into the method. The second term represents the variance of the chosen model and finally the last terms is variance of the error  $\epsilon$ .

To derive this equation, we need to recall that the variance of  $\mathbf{y}$  and  $\epsilon$  are both equal to  $\sigma^2$ . The mean value of  $\epsilon$  is by definition equal to zero. Furthermore, the function  $f$  is not a stochastics variable, idem for  $\tilde{\mathbf{y}}$ . We use a more compact notation in terms of the expectation value

$$\mathbb{E}[(\mathbf{y} - \tilde{\mathbf{y}})^2] = \mathbb{E}[(\mathbf{f} + \epsilon - \tilde{\mathbf{y}})^2],$$

and adding and subtracting  $\mathbb{E}[\tilde{\mathbf{y}}]$  we get

$$\mathbb{E}[(\mathbf{y} - \tilde{\mathbf{y}})^2] = \mathbb{E}[(\mathbf{f} + \epsilon - \tilde{\mathbf{y}} + \mathbb{E}[\tilde{\mathbf{y}}] - \mathbb{E}[\tilde{\mathbf{y}}])^2],$$

which, using the abovementioned expectation values can be rewritten as

$$\mathbb{E}[(\mathbf{y} - \tilde{\mathbf{y}})^2] = \mathbb{E}[(\mathbf{y} - \mathbb{E}[\tilde{\mathbf{y}}])^2] + \text{Var}[\tilde{\mathbf{y}}] + \sigma^2,$$

that is the rewriting in terms of the so-called bias, the variance of the model  $\tilde{\mathbf{y}}$  and the variance of  $\epsilon$ .

## Understanding what happens

### Summing up

The bias-variance tradeoff summarizes the fundamental tension in machine learning, particularly supervised learning, between the complexity of a model and the amount of training data needed to train it. Since data is often limited, in practice it is often useful to use a less-complex model with higher bias, that is a model whose asymptotic performance is worse than another model because it is easier to train and less sensitive to sampling noise arising from having a finite-sized training dataset (smaller variance).

The above equations tell us that in order to minimize the expected test error, we need to select a statistical learning method that simultaneously achieves low variance and low bias. Note that variance is inherently a nonnegative quantity, and squared bias is also nonnegative. Hence, we see that the expected test MSE can never lie below  $\text{Var}(\epsilon)$ , the irreducible error.

What do we mean by the variance and bias of a statistical learning method? The variance refers to the amount by which our model would change if we estimated it using a different training data set. Since the training data are used to fit the statistical learning method, different training data sets will result in a different estimate. But ideally the estimate for our model should not vary too much between training sets. However, if a method has high variance then small changes in the training data can result in large changes in the model. In general, more flexible statistical methods have higher variance.

You may also find this recent [article](#) of interest.

## The same example but now with cross-validation

### Cross-validation with Ridge

### Applying Regression Analysis to the Ising Model

#### The Ising model

The one-dimensional Ising model with nearest neighbor interaction, no external field and a constant coupling constant  $J$  is given by

$$H = -J \sum_k^L s_k s_{k+1}, \quad (2)$$

where  $s_i \in \{-1, 1\}$  and  $s_{N+1} = s_1$ . The number of spins in the system is determined by  $L$ . For the one-dimensional system there is no phase transition.

We will look at a system of  $L = 40$  spins with a coupling constant of  $J = 1$ . To get enough training data we will generate 10000 states with their respective energies.

Here we use ordinary least squares regression to predict the energy for the nearest neighbor one-dimensional Ising model on a ring, i.e., the endpoints wrap around. We will use linear regression to fit a value for the coupling constant to achieve this.

### Reformulating the problem to suit regression

A more general form for the one-dimensional Ising model is

$$H = - \sum_j^L \sum_k^L s_j s_k J_{jk}. \quad (3)$$

Here we allow for interactions beyond the nearest neighbors and a state dependent coupling constant. This latter expression can be formulated as a matrix-product

$$\mathbf{H} = \mathbf{XJ}, \quad (4)$$

where  $X_{jk} = s_j s_k$  and  $J$  is a matrix which consists of the elements  $-J_{jk}$ . This form of writing the energy fits perfectly with the form utilized in linear regression, that is

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}, \quad (5)$$

We split the data in training and test data as discussed in the previous example

## Linear regression

In the ordinary least squares method we choose the cost function

$$C(\mathbf{X}, \boldsymbol{\beta}) = \frac{1}{n} \{(\mathbf{X}\boldsymbol{\beta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\beta} - \mathbf{y})\}. \quad (6)$$

We then find the extremal point of  $C$  by taking the derivative with respect to  $\boldsymbol{\beta}$  as discussed above. This yields the expression for  $\boldsymbol{\beta}$  to be

$$\boldsymbol{\beta} = \frac{\mathbf{X}^T \mathbf{y}}{\mathbf{X}^T \mathbf{X}},$$

which immediately imposes some requirements on  $\mathbf{X}$  as there must exist an inverse of  $\mathbf{X}^T \mathbf{X}$ . If the expression we are modeling contains an intercept, i.e., a constant term, we must make sure that the first column of  $\mathbf{X}$  consists of 1. We do this here

## Singular Value decomposition

Doing the inversion directly turns out to be a bad idea since the matrix  $\mathbf{X}^T \mathbf{X}$  is singular. An alternative approach is to use the **singular value decomposition**. Using the definition of the Moore-Penrose pseudoinverse we can write the equation for  $\boldsymbol{\beta}$  as

$$\boldsymbol{\beta} = \mathbf{X}^+ \mathbf{y},$$

where the pseudoinverse of  $\mathbf{X}$  is given by

$$\mathbf{X}^+ = \frac{\mathbf{X}^T}{\mathbf{X}^T \mathbf{X}}.$$

Using singular value decomposition we can decompose the matrix  $\mathbf{X} = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^T$ , where  $\mathbf{U}$  and  $\mathbf{V}$  are orthogonal(unitary) matrices and  $\boldsymbol{\Sigma}$  contains the singular values (more details below). where  $\mathbf{X}^+ = \mathbf{V}\boldsymbol{\Sigma}^+ \mathbf{U}^T$ . This reduces the equation for  $\boldsymbol{\beta}$  to

$$\boldsymbol{\beta} = \mathbf{V}\boldsymbol{\Sigma}^+ \mathbf{U}^T \mathbf{y}. \quad (7)$$

Note that solving this equation by actually doing the pseudoinverse (which is what we will do) is not a good idea as this operation scales as  $\mathcal{O}(n^3)$ , where  $n$  is the number of elements in a general matrix. Instead, doing  $QR$ -factorization and solving the linear system as an equation would reduce this down to  $\mathcal{O}(n^2)$  operations.

When extracting the  $J$ -matrix we need to make sure that we remove the intercept, as is done here

A way of looking at the coefficients in  $J$  is to plot the matrices as images.

It is interesting to note that OLS considers both  $J_{j,j+1} = -0.5$  and  $J_{j,j-1} = -0.5$  as valid matrix elements for  $J$ . In our discussion below on hyperparameters and Ridge and Lasso regression we will see that this problem can be removed, partly and only with Lasso regression.

In this case our matrix inversion was actually possible. The obvious question now is what is the mathematics behind the SVD?

## The one-dimensional Ising model

Let us bring back the Ising model again, but now with an additional focus on Ridge and Lasso regression as well. We repeat some of the basic parts of the Ising model and the setup of the training and test data. The one-dimensional Ising model with nearest neighbor interaction, no external field and a constant coupling constant  $J$  is given by

$$H = -J \sum_k^L s_k s_{k+1}, \quad (8)$$

where  $s_i \in \{-1, 1\}$  and  $s_{N+1} = s_1$ . The number of spins in the system is determined by  $L$ . For the one-dimensional system there is no phase transition.

We will look at a system of  $L = 40$  spins with a coupling constant of  $J = 1$ . To get enough training data we will generate 10000 states with their respective energies.

A more general form for the one-dimensional Ising model is

$$H = - \sum_j^L \sum_k^L s_j s_k J_{jk}. \quad (9)$$

Here we allow for interactions beyond the nearest neighbors and a more adaptive coupling matrix. This latter expression can be formulated as a matrix-product on the form

$$H = XJ, \quad (10)$$

where  $X_{jk} = s_j s_k$  and  $J$  is the matrix consisting of the elements  $-J_{jk}$ . This form of writing the energy fits perfectly with the form utilized in linear regression,

viz.

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}. \quad (11)$$

We organize the data as we did above

We will do all fitting with **Scikit-Learn**,

When extracting the  $J$ -matrix we make sure to remove the intercept And then we plot the results The results perfectly with our previous discussion where we used our own code.

## Ridge regression

Having explored the ordinary least squares we move on to ridge regression. In ridge regression we include a **regularizer**. This involves a new cost function which leads to a new estimate for the weights  $\boldsymbol{\beta}$ . This results in a penalized regression problem. The cost function is given by

$$C(\mathbf{X}, \boldsymbol{\beta}; \lambda) = (\mathbf{X}\boldsymbol{\beta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\beta} - \mathbf{y}) + \lambda \boldsymbol{\beta}^T \boldsymbol{\beta}. \quad (12)$$

## LASSO regression

In the **Least Absolute Shrinkage and Selection Operator** (LASSO)-method we get a third cost function.

$$C(\mathbf{X}, \boldsymbol{\beta}; \lambda) = (\mathbf{X}\boldsymbol{\beta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\beta} - \mathbf{y}) + \lambda \sqrt{\boldsymbol{\beta}^T \boldsymbol{\beta}}. \quad (13)$$

Finding the extremal point of this cost function is not so straight-forward as in least squares and ridge. We will therefore rely solely on the function “Lasso“ from **Scikit-Learn**.

It is quite striking how LASSO breaks the symmetry of the coupling constant as opposed to ridge and OLS. We get a sparse solution with  $J_{j,j+1} = -1$ .

## Performance as function of the regularization parameter

We see how the different models perform for a different set of values for  $\lambda$ .

We see that LASSO reaches a good solution for low values of  $\lambda$ , but will "wither" when we increase  $\lambda$  too much. Ridge is more stable over a larger range of values for  $\lambda$ , but eventually also fades away.

## Finding the optimal value of $\lambda$

To determine which value of  $\lambda$  is best we plot the accuracy of the models when predicting the training and the testing set. We expect the accuracy of the training set to be quite good, but if the accuracy of the testing set is much lower this tells us that we might be subject to an overfit model. The ideal scenario is an accuracy on the testing set that is close to the accuracy of the training set.



From the above figure we can see that LASSO with  $\lambda = 10^{-2}$  achieves a very good accuracy on the test set. This by far surpasses the other models for all values of  $\lambda$ .